

```
In [1]: import tensorflow as tf
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Flatten
        from tensorflow.keras.datasets import mnist
        from tensorflow.keras.utils import to_categorical
```

```
In [2]: # Load the MNIST dataset
        (x_train, y_train), (x_test, y_test) = mnist.load_data()

        # Normalize pixel values to [0, 1]
        x_train, x_test = x_train / 255.0, x_test / 255.0

        # One-hot encode the labels
        y_train = to_categorical(y_train, 10)
        y_test = to_categorical(y_test, 10)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 ————— 0s 0us/step

```
In [3]: # Define the model
        model = Sequential([
            Flatten(input_shape=(28, 28)),      # Flatten 28x28 input images into 1D vector
            Dense(128, activation='relu'),      # Hidden Layer with 128 neurons
            Dense(10, activation='softmax')     # Output layer with 10 neurons (one per class)
        ])
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

```
In [4]: # Compile the model
        model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [5]: # Train the model
        history = model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.2)
```

Epoch 1/5

1500/1500 ————— 6s 3ms/step - accuracy: 0.8622 - loss: 0.4747 - val_accuracy: 0.9548 - val_loss: 0.1599

Epoch 2/5

1500/1500 ————— 7s 2ms/step - accuracy: 0.9599 - loss: 0.1352 - val_accuracy: 0.9673 - val_loss: 0.1138

Epoch 3/5

1500/1500 ————— 5s 2ms/step - accuracy: 0.9738 - loss: 0.0903 - val_accuracy: 0.9722 - val_loss: 0.0954

Epoch 4/5

1500/1500 ————— 3s 2ms/step - accuracy: 0.9804 - loss: 0.0634 - val_accuracy: 0.9732 - val_loss: 0.0912

Epoch 5/5

1500/1500 ————— 5s 2ms/step - accuracy: 0.9863 - loss: 0.0476 - val_accuracy: 0.9696 - val_loss: 0.0989

```
In [6]: # Evaluate the model on test data
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Accuracy: {test_accuracy * 100:.2f}%')
```

313/313 ————— 1s 3ms/step - accuracy: 0.9709 - loss: 0.0977
Test Accuracy: 97.46%